

1	M101J - MongoDB for Java Developers .....	2
1.1	General .....	2
1.2	Week 2 – CRUD .....	2
1.2.1	Insert, Find & Count .....	2
1.2.2	Updates .....	2
1.2.3	Deletes .....	3
1.2.4	Get last error .....	3
1.2.5	JavaDriver .....	3
1.3	Week 3 – Schema Design .....	4
1.3.1	One-to-one relations .....	4
1.3.2	One-to-many relations .....	4
1.3.3	Many-to-many relations .....	4
1.3.4	GridFS - Blobs .....	4
1.4	Week 4 – Performance .....	5
1.4.1	Index creation and deletion .....	5
1.4.2	Multi key indices .....	5
1.4.3	Sparse Index .....	5
1.4.4	Explain & Hint .....	5
1.4.5	Efficiency of indexes .....	5
1.4.6	Geospatial indexes .....	6
1.4.7	Profiling slow queries .....	6
1.4.8	mongotop & mongostat .....	6
1.5	Week 5 – Aggregation .....	6
1.5.1	Aggregation Pipeline .....	6
1.5.2	Compound grouping .....	7
1.5.3	Group stage .....	7
1.5.4	Project stage .....	7
1.5.5	Match stage .....	7
1.5.6	Sort stage .....	7
1.5.7	Skip and limit stage .....	7
1.5.8	Unwind stage .....	7
1.5.9	Limitations of the aggregation framework .....	7
1.6	Week 6 - Replication .....	7
1.6.1	Failover and Rollback .....	8
1.6.2	Connecting from the Java Driver .....	8
1.6.3	Write Concerns .....	8
1.6.4	Read Preferences .....	9
1.7	Week 6 – Sharding .....	9
1.7.1	Building a sharded environment .....	9
1.7.2	Sharding implications .....	10
1.7.3	Sharding key .....	10

# 1 M101J - MongoDB for Java Developers

## 1.1 General

- `mongoimport -d myDB -c xxx [--drop] < xxx.json` // imports a JSON file into a collection

## 1.2 Week 2 – CRUD

### 1.2.1 Insert, Find & Count

- `db.[collection].insert({...})`
- The collection unique ID field is called “\_id” and can be provided. If not provided an ObjectId will be generated based on the time, machine, process-id and process dependent counter.
- “\_id” does not have to be a scalar value – it can be a document, e.g. `_id : {a:1, b:'ronald'}`
- `db.[collection].find || findOne ({...}, {field1 : true, ...}).pretty()` //no argument will find all docs
- `db.[collection].count({...})`

#### 1.2.1.1 Operators

- Ranges: `{myField: {$gt: 100, $lt: 10}}` `$gte`, `$lte` → Can be applied to numbers and strings (ASCII)
- Regex: `{myField: {$regex: “a$”}}`
- Set operators: `{myField: {$in: [“one”, “two”, ...]}}` `$nin`
- Boolean:
  - `{ $or: [{...}, {...}, ...]}` `$and`
  - `$not` - negates result of other operation or regular expression query
  - tags: `{ $ne: “gardening”}` // works on keys pointing to single values or arrays – inefficient – can’t use indexes
  - `{myField: { $exists: true}}` //checks if particular key exists in document
- `{myField: { $type: 2}}` // 2=String as defined in BSON spec
- <http://docs.mongodb.org/manual/reference/operator/>

#### 1.2.1.2 Array-Operators

- `{myArrayField: “test”}` → Will find any documents where the array contains the value “test”
- `{myArrayField.0 : “test”}` → Value at particular position within array
- `{myArrayField: { $all: [“one”, “two”, ...]}}` → array contains all given values in any order
- `{myArrayField: { $size: 3}}` → array with three elements

#### 1.2.1.3 Nested Documents

- `{“myField.mySubfield” : “test”}` → Dot-Notation needs to be put in “”
- `{“myArrayField.0.mySubfield” : “test”}` → stipulate zeroth element of array
- `{“myArrayField.mySubfield” : “test”}` → search in any of the array elements
- `{myArrayField : { $elemMatch: { mySubfield : “test”, mySubfield2 : “test2”}}}` → restrict multiple conditions to same subdocument of array field

#### 1.2.1.4 Cursors

- `myCursor = db.[collection].find(); null;` → append null as not to print out the cursor immediately
- `myCursor.hasNext()` `myCursor.next()`
- `myCursor.skip(2).limit(5).sort({name : -1}); null;` → modifies the query executed on the server

## 1.2.2 Updates

- `db.[collection].update({ myQuery }, {myField: “newValue”, ... })` → replaces the existing document
- `db.[collection].update({ myQuery }, { $set : {myField: “newValue”}})` → Create or update myField
- `db.[collection].update({ myQuery }, { $inc : {age: 1}})`

- `db.[collection].update({ myQuery }, { $unset : { myField: 1 } })`
- `db.[collection].update({ myQuery }, { $set : { myField: "newValue" }, { upsert: true } })` → Create or update document specified by { myQuery } with myField

#### Arrays

- `db.[collection].update({ myQuery }, { $set : { "myArray.2": "x" } })` → Set 3<sup>rd</sup> position of Array
- `db.[collection].update({ myQuery }, { $push : { myArray: "y" } })`
- `db.[collection].update({ myQuery }, { $addToSet : { myArray: "y" } })` // will only add if does not exist yet
- `db.[collection].update({ myQuery }, { $pop : { myArray: 1 } })` // pop right-most
- `db.[collection].update({ myQuery }, { $pop : { myArray: -1 } })` // pop left-most
- `db.[collection].update({ myQuery }, { $pushAll : { myArray: [ "a", "b", "c" ] } })`
- `db.[collection].update({ myQuery }, { $pull : { myArray: "c" } })` // remove value "c"
- `db.[collection].update({ myQuery }, { $pullAll : { myArray: [ "a", "b", "c" ] } })`

#### Multi-Update

- `db.[collection].update({}, { $set: { title: "Dr." }, { multi: true } })` // {} matches all documents

### 1.2.3 Deletes

- `db.[collection].remove({})` // no argument will remove all documents from collection – not isolated
- `db.[collection].drop()` // faster but drops collection including indexes

### 1.2.4 Get last error

- `db.runCommand({ getLastError: 1 })` // get info/error of last operation on this connection to mongod

### 1.2.5 JavaDriver

```
MongoClient mongo = new MongoClient();
```

```
DB db = client.getDB("test");
```

```
DBCollection collection = db.getCollection("test");
```

```
DBObject document = new BasicDBObject("key", "value").append("key2", "value2");
collection.insert(document);
```

```
DBObject query = new BasicDBObject("key", "value").append("y new BasicDBObject("$gt", 0).append("$lt", 99));
query = new QueryBuilder().start("key").is("value").and("y").greaterThan(0).lessThan(99).get();
document = collection.findOne();
collection.count(query);
```

```
DBObject fieldSelector = new BasicDBObject("key": true).append("_id", false);
DBCursor cursor = collection.find(query, fieldSelector).sort(sortDocument).skip(2).limit(5);
cursor.hasNext(); document = cursor.next();
finally { cursor.close(); }
```

```
collection.update(query, new BasicDBObject("$set", new BasicDBObject("lala", 1)));
collection.remove();
```

```
collection.findAndModify(query, fields, sortCriteria, removeDocumentFlag, updateDocument, returnNewFlag, upsertFlag); // is atomic
```

## 1.3 Week 3 – Schema Design

- Which data is used together; which data is read; which data is written all the time → *Make the schema matching to the data access patterns of your application*
- Mongo has no joins / no foreign key constraints → *but can embed Documents (Pre-Join)*
- MongoDB does not support transactions → *but atomic operations on ONE document, so instead of transactions, you have three options*
  - *Restructure data to live in one document*
  - *Implement transaction in Software*
  - *Tolerate a little bit of inconsistency*

### 1.3.1 One-to-one relations

- EITHER two collections were one document point to the document in the other collection by `_id`
- OR embed one document in the other
- Decision driven by
  - Frequency of access; Are the documents read together – do you want to pull everything into memory
  - Are the documents written together
  - Document max size: 16 MB
  - Do you need atomicity

### 1.3.2 One-to-many relations

- Two collections – linking from “many collection” to `_id` of “one collection”
- If it is really one-to-few: possible to have “one collection” and embed “few document”

### 1.3.3 Many-to-many relations

- It often really is a few-to-few relation, e.g. authors-books
- EITHER have two collections and add an array of book-ids in author document or vice versa – depends on access pattern
- (OR embed book in author document but this might lead to inconsistency as one book might be duplicated – also not a good idea if you need to store one item before the other exists, e.g. student-teacher)

### 1.3.4 GridFS - Blobs

- > 16MB
- Files collection and Chunks collection. MongoDB spits Blob into chunks of 16MB and stores them in the Chunks collection. Each chunk has a `file_id` pointing to the `_id` of its file document.

```
GridFS videos = new GridFS(db, "videos");
GridFSInputFile video = videos.createFile(inputStream, "video.mp4");
BasicDBObject metadata = ...
video.setMetaData(metadata);
video.save();
...
GridFSDBFile myVideo = videos.findOne(new BasicDBObject("filename", "video.mp4"));
myVideo.writeTo(outputStream)
```

## 1.4 Week 4 – Performance

### 1.4.1 Index creation and deletion

- `db.[collection].ensureIndex({student_id:1})`
- `1=ascending, -1=descending` → important for sorting not so much for searching
- → Sorting can also use a reverse index if the sort criteria are exactly the reverse of an (simple or compound) index
- Compound Index:
  - `db.[collection].ensureIndex({student_id:1,class:-1})`
  - General rule: A Query where one term demands an exact match and another specifies a range requires a compound index where the range key comes second
- Unique Index: `db.[collection].ensureIndex({student_id:1}, {unique: true})` // `dropDups: true` → dangerous
- By default index creation is done in the foreground which is fast but blocking all other writers to the same DB. Background index creation `{background: true}` will be slow but it will not block the writers
- We want indexes to be in memory. Find out the index size: `db.[col].stats()` or `db.[col].totalIndexSize()`
- `db.system.indexes.find()` → finds all indexes of the current db
- `db.[collection].getIndexes()` → all indexes of collection
- `db.[collection].dropIndex({student_id:1})`

### 1.4.2 Multi key indices

- A multi key index is an index on an array field of a document, e.g. a student document has array of teacher-ids. One can add a multi key index on the teachers-array, which indexes all of the values in the array for all the documents.
- Multi key indices are one of the reason that linking works so well in MongoDB
- It is not possible to have a compound index with two array (multi key) fields

### 1.4.3 Sparse Index

- Missing index key in documents map to null → unique key not possible because multiple nulls are not allowed
- Sparse indexes only index documents that have a key set for the key being indexed `{unique: true, sparse: true}`
- On a sorted find the non-indexed documents will not be found when the sparse index is used for the sort

### 1.4.4 Explain & Hint

- `db.[collection].find({...}).explain()`
- `db.[collection].find({...}).explain(true)` //shows all possible plans
- `db.[collection].find({...}).hint({a:1, b:1})` // use specified index
- `db.[collection].find({...}).hint({$natural:1})` // use no index
- In Java:
  - `.find(query).hint("IndexName")` OR
  - `.find(query).hint(new BasicDocument(a, 1).append(b, 1))`

### 1.4.5 Efficiency of indexes

- `$gt`, `$lt`, `$ne`, `$nin`, `$not($exists)` might be inefficient even if an index is used because still many index items (indexed documents) need to be scanned → may be a good idea to use a hint to use a diff. index
- `$regex` can only use an index if it is stemmed on the left side, e.g. `/^abc/`

### 1.4.6 Geospatial indexes

- `.ensureIndex({location: '2d', type: 1})` // Compound index on location (uses 2d-index) and ascending type

### 1.4.7 Profiling slow queries

- MongoDB logs slow queries (>100ms) by default into the logfile
- Use profiler:
- `mongod --profile 1 --slowms 10` // logs all queries taking longer than 10ms to `system.profile` collection
- Levels: 0=off (default) 1=log slow queries 2=log all queries (general debugging feature for dev.)
- Mongo shell: `db.getProfilingLevel()` `db.getProfilingStatus()` `db.setProfilingLevel(level, slowms)`
- `mongod --notablescan` option: Set `notablescan = true` on your dev or test machine to find operations that require a table scan

### 1.4.8 mongotop & mongostat

- `mongotop 3` // runs every 3 seconds showing you in which collection how much time (read, write, total) is spent
- `mongostat` // shows inserts, queries, updates, deletes, ... per second
- → `idx miss %` = index which could not be accessed in memory

## 1.5 Week 5 – Aggregation

- group by like-functionality
- `db[col].aggregate([  
 {$group: {  
 _id: '$manufacturer', → set the field '_id' to the field you want to group by  
 count: {$sum: 1} → define field 'count'  
 }}  
])`

### 1.5.1 Aggregation Pipeline

- Each document in the array parameter to the aggregate function is a stage in the pipeline
- E.g. Collection → project stage → match stage → group stage → sort stage → result
- Stages:
  - `$project`: Select relevant fields and reshape document (in: 1 / out: 1)
  - `$match`: Filters documents; (in: n / out: n-x)
  - `$group`: Aggregates; Reduces the number of documents (in: n / out: n-x)
  - `$sort`: Sorts the documents (in: 1 / out: 1)
  - `$skip`: Skips documents (in: n / out: n-x)
  - `$limit`: Limits returned documents (in: n / out: n-x)
  - `$unwind`: Explodes arrays - Produces a document for each value in an array-key-field with everything else repeated (in: n / out: n+x)
- Each stage can exist more than once in a pipeline

Example:

```
db.zips.aggregate([  
  {$match: {state: {$in: ["CA", "NY"]}}},  
  {$group: {_id: {city: "$city", state: "$state"}, pop: {$sum: "$pop"}}},  
  {$match: {pop: {$gt: 25000}}},  
  {$group: {_id: null, avg: {$avg: "$pop"}}}  
])
```

### 1.5.2 Compound grouping

- Use a compound id: `_id: {myManufacturer: "$manufacturer", myCategory: "$category"}, ...`

### 1.5.3 Group stage

- `$sum`: Add one to a key (→ `mySum: {$sum:1}`) or sum up keys (→ `sum_prices:{$sum:"$price"}`)
- `$avg`, `$min`, `$max`: Average, Minimum or maximum value of a key
- Create arrays: `$push`, `$addToSet` → `categories: {$addToSet: "$category"}`
- Only useful after a sort: `$first`, `$last` → `{$group:{_id:"$_id.state", population:{$first:"$population"}}`

### 1.5.4 Project stage

- Remove keys - If you don't mention a key, it is not included, except for `_id`, which must be explicitly suppressed `{$project: {_id: 0, ...`
- Add keys (also possible to create new subdocuments)
- Keep keys: `{$project: {myKey: 1, ...`
- Rename keys / Use functions: `$toUpper`, `$toLower`, `$add`, `$multiply`

### 1.5.5 Match stage

- `{$match: {pop: {$gt: 100000}}}`

### 1.5.6 Sort stage

- Memory intensive; Can't use index (at least after grouping)
- `{$sort: {population: -1}}`

### 1.5.7 Skip and limit stage

- Makes only sense when you do a sort first
- First skip – then limit (order of the stages in the pipeline matter)

### 1.5.8 Unwind stage

- `{$unwind: "$tags"}`

### 1.5.9 Limitations of the aggregation framework

- A result document can only be **16GB**
- One can only use **up to 10% of the memory** on a machine
- In sharded environment: After first `$group` / `$sort` the next phase have to be performed on the mongos router
- Alternative to aggregation framework: map-reduce

SQL Comparison: <http://docs.mongodb.org/manual/reference/sql-aggregation-comparison/>

## 1.6 Week 6 - Replication

- You can only write to the **primary node** which replicates asynchronous to **secondary nodes**
- If you only read from the primary you will have strong consistency (default behaviour)
- You can allow your reads to go to secondaries → you might read stale data and have eventual consistency
- If the primary goes down the secondaries elect a new primary and the Java driver automatically connects to the new primary
- **Arbiter nodes** exist for voting purposes, e.g. if you have an even number of regular (= primary & secondary) nodes it can ensure a majority in an election → Allows you to have only two regular nodes.

- **Delayed nodes** are disaster recovery nodes: Can be set to be whatever time behind the regular nodes. Can't participate in elections
- **Hidden node** (e.g. for reporting) can't become the primary but can participate in elections
- Start a replication set: `mongod -replSet m101 --logpath "1.log" --dbpath /data/rs1 -fork`
- Register replica set nodes in the mongo shell:

```
config = { _id: "m101", members:[
    { _id : 0, host : "localhost:27017", priority:0, slaveDelay:5},
    { _id : 1, host : "localhost:27018"},
    { _id : 2, host : "localhost:27019"} ]
}
```

**rs.initiate(config)** ← Initializes the replica set // can't be executed on a node which can't become primary

**rs.status()** ← Gives you the status information about the replica set

**rs.isMaster()** ← Tells you if you are the primary node

**rs.slaveOk()** ← If issued on a secondary node it allows you to read from this secondary node

**rs.stepDown()** ← Forces primary node to step down as a primary node

- Replication is done via a capped collection called **oplog.rs** in the "local" database
- Secondaries ask the primary for any items since a certain timestamp

### 1.6.1 Failover and Rollback

Scenario:

- When the primary dies a secondary which becomes elected as a new primary which does not have the latest entries from the old primaries oplog
- When the former primary node comes back up as a secondary node it will request the oplog data from the new primary and roll back the writes the current primary does not have and write them to a rollback file which can be applied manually
- If the oplog of the new primary has looped during the time the old primary was down the entire dataset will be copied from the new primary
- The risk of losing data due to a rollback can be avoided by waiting till the majority of the nodes have the data → set the write concern `w=majority`

### 1.6.2 Connecting from the Java Driver

Provide a **seed list** to the MongoClient instance

```
new MongoClient(Arrays.asList(
    new ServerAddress("localhost", 27017),
    new ServerAddress("localhost", 27018),
    new ServerAddress("localhost", 27019),
));
```

→ Will work even if the primary is not part of the seed list. The Java Client starts a background thread which pings all nodes from the seed list and all discovered nodes to find out which one is the primary

### 1.6.3 Write Concerns

Client writes to a primary:

1. Primary writes into RAM (collection and oplog)
  2. The writes are asynchronously journaled (Gives recoverability in case of a crash)
  3. The writes are written into the data directory
  4. Secondaries are replicating writes from the primary's oplog
- The insert method sent from the Client does not expect a response
  - The client sends a second command "getLastError"



- `w=0` → Unacknowledged; Fast writes – no “getLastError” command
- `w=1` → Will wait for the primary to write into RAM
- `w=[n]` → Will wait for the primary and n-1 to write into RAM
- `w=majority` → Wait for majority of replica set to write to RAM
- `j=true` → Will wait for the primary to write into the journal
- `fsync=true` → Will wait for the primary to write into the data directory
- `wtimeout=[milliseconds]` → Indicate how long you are willing to wait
- Java Driver
  - Default: `w=1` and `wtimeout=0` (=infinite)
  - `client | db | collection.setWriteConcern(WriteConcern.JOURNALED)`
  - `collection.insert(doc, WriteConcern.JOURNALED)`
  - `WriteConcern.JOURNALED = new WriteConcern([w=]1,[wtimeout=]0, [fsync=]false,[j=]true)`

#### 1.6.4 Read Preferences

- Primary → All reads are sent to the primary (default to guarantee strict consistency)
- Secondary → Send reads to randomly selected secondaries, but not to the primary
- Secondary Preferred → Send reads to secondaries or to the primary if all secondaries are down
- Primary Preferred → Sends reads to primary or to a secondary if primary is down
- Nearest → Send reads to secondary or primary
- If you read from secondaries you might get stale reads. The might be OK if different applications do the writes and the reads
- For any read preference (except Primary) the driver will look at ping times and will only send reads to nodes which are within the latency window (15ms) of the fastest
- `client | db | collection.setReadPreference(ReadPreference.primaryPreferred())`
- `collection.find().setReadPreference(ReadPreference.nearest());`

### 1.7 Week 6 – Sharding

- Enables horizontal scalability
- Shards are typically itself replica sets
- `mongos` is the sharding router which distributes data to the individual shards
- The application (and also the mongo shell) connects to `mongos` instead of `mongod`
- There can be multiple `mongos` - `mongos` typically runs on the same server as the application
- If a `mongos` goes down the application will connect to a different one – similar to replica sets
- Shard key determines to which shard a document goes
- Sharding is at database level – but you can define if you want to shared or not shard a specific collection
- `Config servers` (which are `mongod`) keep track of where the shards are – in production you typically use 3 of them

#### 1.7.1 Building a sharded environment

Set up two shards each a replica set of three `mongod` nodes

- Set up shard as a replication set:
- `mongod --replSet s0 --logpath "s0-r0.log" --dbpath /data/shard0/rs0 --port 37017 --fork --shardsvr`
- Set up config server:
- `mongod --logpath "cfg-a.log" --dbpath /data/config/config-a --port 57040 --fork --configsvr`
- Set up `mongos` router with information about the config servers:
- `mongos --logpath "mongos-1.log" --configdb localhost:57040,localhost:57041,localhost:57042 --fork`
- `mongos` now listens on the default `mongod` port
- On the mongo shell – tell the config servers (via the `mongos`) about the shards:
- `db.adminCommand( { addshard : "s0/"+"localhost:37017" } );`

- ... add further shards
- `db.adminCommand({enableSharding: "test"})` → enable sharding on test DB
- `db.adminCommand({shardCollection: "test.grades", key: {student_id:1}});` → shard collection "grades" with the shard key "student\_id"
- `sh.help()` → Will display all the shard commands available in the mongo shell, e.g. `sh.status()`

### 1.7.2 Sharding implications

- Needs an index on first element of the shard key (can be compound but not multi-key)
- Each document needs the shard key
- The shard key is immutable
- On an update you need to specify the shard key or specify multi: true
- A find with no shard key will go to all shards
- The key used in most queries should be the shard key
- You can't have a unique index unless it is part of/starts with the shard key
- Write concerns are still important in a sharded setup

### 1.7.3 Sharding key

- Sufficient cardinality (variety of values)
- Avoid monotonically increasing keys to avoid hotspotting in writing (e.g. order\_id, order\_date)
- Compound sharding key is possible

